

Experiment No. 2
Analyze the Titanic Survival Dataset and apply appropriate regression technique
Date of Performance:
Date of Submission:

**Aim:** Analyze the Titanic Survival Dataset and apply appropriate Regression Technique.

**Objective:** Able to perform various feature engineering tasks, apply logistic regression on the given dataset and maximize the accuracy.

### Theory:


Logistic Regression was used in the biological sciences in early twentieth century. It was then used in many social science applications. Logistic Regression is used when the dependent variable(target) is categorical and is binary in nature. In order to perform binary classification the logistic regression techniques makes use of Sigmoid function.

For example,

To predict whether an email is spam (1) or (0)

Whether the tumor is malignant (1) or not (0)

Consider a scenario where we need to classify whether an email is spam or not. If we use linear regression for this problem, there is a need for setting up a threshold based on which classification can be done. Say if the actual class is malignant, predicted continuous value 0.4 and the threshold value is 0.5, the data point will be classified as not malignant which can lead to serious consequence in real time.



From this example, it can be inferred that linear regression is not suitable for classification problem. Linear regression is unbounded, and this brings logistic regression into picture. Their value strictly ranges from 0 to 1.

## Dataset:

The sinking of the Titanic is one of the most infamous shipwrecks in history.

On April 15, 1912, during her maiden voyage, the widely considered “unsinkable” RMS Titanic sank after colliding with an iceberg. Unfortunately, there weren’t enough lifeboats for everyone onboard, resulting in the death of 1502 out of 2224 passengers and crew.

While there was some element of luck involved in surviving, it seems some groups of people were more likely to survive than others.

In this challenge, we ask you to build a predictive model that answers the question: “what sorts of people were more likely to survive?” using passenger data (ie name, age, gender, socio-economic class, etc).

Variable	Definition	Key
survival	Survival	0 = No, 1 = Yes
pclass	Ticket class	1 = 1st, 2 = 2nd, 3 = 3rd
sex	Sex	
Age	Age in years	
sibsp	# of siblings / spouses aboard the Titanic	
parch	# of parents / children aboard the Titanic	
ticket	Ticket number	
fare	Passenger fare	
cabin	Cabin number	
embarked	Port of Embarkation	C = Cherbourg, Q = Queenstown, S = Southampton

## Variable Notes

pclass: A proxy for socio-economic status (SES)

1st = Upper, 2nd = Middle, 3rd = Lower

age: Age is fractional if less than 1. If the age is estimated, is it in the form of xx.5

sibsp: The dataset defines family relations in this way...,

Sibling = brother, sister, stepbrother, stepsister

Spouse = husband, wife (mistresses and fiancés were ignored)

parch: The dataset defines family relations in this way...

Parent = mother, father

Child = daughter, son, stepdaughter, stepson

Some children travelled only with a nanny, therefore parch=0 for them.

### **Code:**

### **Conclusion:**

1. What are features have been chosen to develop the model? Justify the features chosen to determine the survival of a passenger.
  - Firstly the Titanic dataset was loaded and preprocessing was carried out in order to clean the data. Unwanted columns like alive, alone, embarked\_town, who, adult\_male & deck was dropped as they were of no use for prediction the model. Get\_dummies function was used in order to convert categorical data into numeric data so the prediction can be done easily.
  - Different diagrams are plotted for the visualization and counts of number of female and males are done and also null vales are dropped from embarked column and mean values are filled in place of null values in column age.
  - Further to determine survival of a passenger the data was split into train and test data with size of 80 & 20 where Y\_train contains column = Survival and X\_Train consists of the remaining columns which will help us to predict the outcome.

**2. Comment on the accuracy obtained.**

Accuracy is the percentage of correct classifications that a trained machine learning model achieves, i.e., the number of correct predictions divided by the total number of predictions across all classes. The model is trained using Logistic Regression and the accuracy obtained is  $= 0.84$  which is equivalent to approximately 84%.

```
import seaborn as sns
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
df = sns.load_dataset("titanic")
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 15 columns):
#   Column      Non-Null Count  Dtype
---  -
0   survived    891 non-null    int64
1   pclass      891 non-null    int64
2   sex         891 non-null    object
3   age         714 non-null    float64
4   sibsp       891 non-null    int64
5   parch       891 non-null    int64
6   fare        891 non-null    float64
7   embarked    889 non-null    object
8   class       891 non-null    category
9   who         891 non-null    object
10  adult_male  891 non-null    bool
11  deck        203 non-null    category
12  embark_town 889 non-null    object
13  alive       891 non-null    object
14  alone       891 non-null    bool
dtypes: bool(2), category(2), float64(2), int64(4), object(5)
memory usage: 80.7+ KB
```

```
df.head()
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	ac
0	0	3	male	22.0	1	0	7.2500	S	Third	man	
1	1	1	female	38.0	1	0	71.2833	C	First	woman	
2	1	3	female	26.0	0	0	7.9250	S	Third	woman	
3	1	1	female	35.0	1	0	53.1000	S	First	woman	
4	0	3	male	35.0	0	0	8.0500	S	Third	man	





```
columns = ['alive', 'alone', 'embark_town', 'who', 'adult_male', 'deck']
data1 = df.drop(columns, axis=1)
```

```
data1.head()
```

survived pclass sex age sibsp parch fare embarked class 

data1

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	
0	0	3	male	22.0	1	0	7.2500	S	Third	
1	1	1	female	38.0	1	0	71.2833	C	First	
2	1	3	female	26.0	0	0	7.9250	S	Third	
3	1	1	female	35.0	1	0	53.1000	S	First	
4	0	3	male	35.0	0	0	8.0500	S	Third	
...	...	...	...	...	...	...	...	...	...	
886	0	2	male	27.0	0	0	13.0000	S	Second	
887	1	1	female	19.0	0	0	30.0000	S	First	
888	0	3	female	NaN	1	2	23.4500	S	Third	
889	1	1	male	26.0	0	0	30.0000	C	First	
890	0	3	male	32.0	0	0	7.7500	Q	Third	

891 rows × 9 columns

```
data1[data1['sex'].str.match("female")].count()
```

```
survived    314
pclass      314
sex         314
age         261
sibsp       314
parch       314
fare        314
embarked    312
class       314
dtype: int64
```

```
data1[data1['sex'].str.match("male")].count()
```

```
survived    577
pclass      577
sex         577
age         453
sibsp       577
parch       577
fare        577
embarked    577
class       577
dtype: int64
```


```
gender = pd.get_dummies(data1['sex'], drop_first=True)
```

```
data1['gender'] = gender
```

```
data1.drop('sex', axis=1,inplace=True)
```


```
sns.countplot(x='survived', hue='pclass', data=data1)
```

<Axes: xlabel='survived', ylabel='count'>



```
plt.hist(data1['age'],bins=30,color='darkred',alpha=0.7)
```

```
(array([30., 14., 10., 14., 3., 29., 39., 65., 73., 42., 65., 62., 33.,
        47., 38., 21., 28., 21., 16., 14., 15., 7., 9., 8., 4., 0.,
        5., 1., 0., 1.]),
array([ 0.42      ,  3.07266667,  5.72533333,  8.378      , 11.03066667,
        13.68333333, 16.336      , 18.98866667, 21.64133333, 24.294      ,
        26.94666667, 29.59933333, 32.252      , 34.90466667, 37.55733333,
        40.21      , 42.86266667, 45.51533333, 48.168      , 50.82066667,
        53.47333333, 56.126      , 58.77866667, 61.43133333, 64.084      ,
        66.73666667, 69.38933333, 72.042      , 74.69466667, 77.34733333,
        80.      ]),
<BarContainer object of 30 artists>)
```





data1.head()

	survived	pclass	age	sibsp	parch	fare	embarked	class	gender
0	0	3	22.0	1	0	7.2500	S	Third	1
1	1	1	38.0	1	0	71.2833	C	First	0
2	1	3	26.0	0	0	7.9250	S	Third	0
3	1	1	35.0	1	0	53.1000	S	First	0
4	0	3	35.0	0	0	8.0500	S	Third	1

```
change = {'First':1 , 'Second':2, 'Third':3}
data1['class'] = data1['class'].replace(change)
```

```
change1 = {'C':1 , 'Q':2, 'S':3}
data1['embarked'] = data1['embarked'].replace(change1)
```

data1.head()

	survived	pclass	age	sibsp	parch	fare	embarked	class	gender
0	0	3	22.0	1	0	7.2500	3.0	3	1
1	1	1	38.0	1	0	71.2833	1.0	1	0
2	1	3	26.0	0	0	7.9250	3.0	3	0
3	1	1	35.0	1	0	53.1000	3.0	1	0
4	0	3	35.0	0	0	8.0500	3.0	3	1

```
column_name = 'embarked'
data1 = data1.dropna(subset = [column_name],axis = 0)
```



```
data1['age'].fillna(data1['age'].mean() , inplace=True)
```

<ipython-input-140-b5d6b2d9217e>:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html)  
data1['age'].fillna(data1['age'].mean() , inplace=True)

```
x=data1.iloc[:,1:]
y=data1.iloc[:,0]
```

x

	pclass	age	sibsp	parch	fare	embarked	class	gender	
0	3	22.000000	1	0	7.2500	3.0	3	1	
1	1	38.000000	1	0	71.2833	1.0	1	0	
2	3	26.000000	0	0	7.9250	3.0	3	0	
3	1	35.000000	1	0	53.1000	3.0	1	0	
4	3	35.000000	0	0	8.0500	3.0	3	1	
...	...	...	...	...	...	...	...	...	
886	2	27.000000	0	0	13.0000	3.0	2	1	

y

```
0      0
1      1
2      1
3      1
4      0
..
886    0
887    1
888    0
889    1
890    0
```

Name: survived, Length: 889, dtype: int64

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
X_train, X_test, Y_train , Y_test = train_test_split(x , y, test_size = 0.2 , random_state=1)
```

```
model = LogisticRegression()
```

```
print(X_train.shape , Y_train.shape)
```

```
(711, 8) (711,)
```

```
model.fit(X_train , Y_train)
```

```
▼ LogisticRegression
LogisticRegression()
```

```
y_pred = model.predict(X_test)
```

```
accuracy = accuracy_score(Y_test,y_pred)
print(f"Accuracy:{accuracy:.2f}")
```

```
Accuracy:0.84
```

✓ 0s completed at 1:21 AM

