



Experiment No. 5
Apply appropriate Unsupervised Learning Technique on the Wholesale Customers Dataset
Date of Performance:
Date of Submission:



Aim: Apply appropriate Unsupervised Learning Technique on the Wholesale Customers Dataset.


Objective: Able to perform various feature engineering tasks, apply Clustering Algorithm on the given dataset.

Theory:

It is basically a type of unsupervised learning method. An unsupervised learning method is a method in which we draw references from datasets consisting of input data without labeled responses. Generally, it is used as a process to find meaningful structure, explanatory underlying processes, generative features, and groupings inherent in a set of examples.

Clustering is the task of dividing the population or data points into a number of groups such that data points in the same groups are more similar to other data points in the same group and dissimilar to the data points in other groups. It is basically a collection of objects on the basis of similarity and dissimilarity between them.

For example: The data points in the graph below clustered together can be classified into one single group. We can distinguish the clusters, and we can identify that there are 3 clusters in the below picture.





Dataset:

This data set refers to clients of a wholesale distributor. It includes the annual spending in monetary units (m.u.) on diverse product categories. The wholesale distributor operating in different regions of Portugal has information on annual spending of several items in their stores across different regions and channels. The dataset consist of 440 large retailers annual spending on 6 different varieties of product in 3 different regions (lisbon , oporto, other) and across different sales channel (Hotel, channel)

Detailed overview of dataset

Records in the dataset = 440 ROWS

Columns in the dataset = 8 COLUMNS

FRESH: annual spending (m.u.) on fresh products (Continuous)

MILK:- annual spending (m.u.) on milk products (Continuous)

GROCERY:- annual spending (m.u.) on grocery products (Continuous)

FROZEN:- annual spending (m.u.) on frozen products (Continuous)

DETERGENTS_PAPER :- annual spending (m.u.) on detergents and paper products (Continuous)

DELICATESSEN:- annual spending (m.u.)on and delicatessen products (Continuous);

CHANNEL: - sales channel Hotel and Retailer

REGION:- three regions (Lisbon, Oporto, Other)

Code:



Conclusion:

Based on the visualization, comment on following:

1. How can you can make use of the clustered data?
 - Clustered data helps to understand different customer segments based on their purchasing behaviour. It will help to design marketing strategies that are more relevant to cluster's preferences.
 - Identifying which products are frequently purchased together within each cluster, we can make personalized product recommendations to customers.
 - Clustering plays an important role in inventory management as it ensures that the right products available in right quantities to meet each cluster's requirements.
 - We can optimize supply chain operations by managing delivery schedules and routes to each cluster's needs.
 - Clustering can also help identify new markets schemes or customer segments that are similar to existing clusters.
2. How the different groups of customers, the *customer segments*, may be affected differently by a specific delivery scheme?
 - If a delivery scheme offers premium or high expenditure delivery options, high-value customers who are willing to pay more may respond positively.
 - Customers who loves shopping frequently may benefit from subscription-based or loyalty-based delivery schemes. These schemes can encourage them to repeat their purchases and loyalty.
 - Some customers may be more price-sensitive and may prefer a cost-effective or free standard delivery scheme.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
df = pd.read_csv('customers.csv')
print(df)
```

	Channel	Region	Fresh	Milk	Grocery	Frozen	Detergents_Paper	\
0	2	3	12669	9656	7561	214	2674	
1	2	3	7057	9810	9568	1762	3293	
2	2	3	6353	8808	7684	2405	3516	
3	1	3	13265	1196	4221	6404	507	
4	2	3	22615	5410	7198	3915	1777	
..	
435	1	3	29703	12051	16027	13135	182	
436	1	3	39228	1431	764	4510	93	
437	2	3	14531	15488	30243	437	14841	
438	1	3	10290	1981	2232	1038	168	
439	1	3	2787	1698	2510	65	477	

	Delicatessen
0	1338
1	1776
2	7844
3	1788
4	5185
..	...
435	2204
436	2346
437	1867
438	2125
439	52

[440 rows x 8 columns]

```
df.head()
```

	Channel	Region	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicatessen
0	2	3	12669	9656	7561	214	2674	1338
1	2	3	7057	9810	9568	1762	3293	1776
2	2	3	6353	8808	7684	2405	3516	7844
3	1	3	13265	1196	4221	6404	507	1788
4	2	3	22615	5410	7198	3915	1777	5185

```
print("Data Types")
df.dtypes
```

```
Data Types
Channel      int64
Region      int64
Fresh       int64
Milk        int64
Grocery     int64
Frozen      int64
Detergents_Paper  int64
Delicatessen int64
dtype: object
```

```
print("Missing values per column:")
print(df.isnull().sum())
```

```
Missing values per column:
Channel      0
Region      0
Fresh       0
Milk        0
Grocery     0
Frozen      0
Detergents_Paper  0
Delicatessen  0
dtype: int64
```

```
print("Descriptive Statistics:")
print(df.describe())
print("Number of duplicate rows: ", df.duplicated().sum())
```

	Channel	Region	Fresh	Milk	Grocery	\
count	440.000000	440.000000	440.000000	440.000000	440.000000	

mean	1.322727	2.543182	12000.297727	5796.265909	7951.277273
std	0.468052	0.774272	12647.328865	7380.377175	9503.162829
min	1.000000	1.000000	3.000000	55.000000	3.000000
25%	1.000000	2.000000	3127.750000	1533.000000	2153.000000
50%	1.000000	3.000000	8504.000000	3627.000000	4755.500000
75%	2.000000	3.000000	16933.750000	7190.250000	10655.750000
max	2.000000	3.000000	112151.000000	73498.000000	92780.000000

	Frozen	Detergents_Paper	Delicatessen
count	440.000000	440.000000	440.000000
mean	3071.931818	2881.493182	1524.870455
std	4854.673333	4767.854448	2820.105937
min	25.000000	3.000000	3.000000
25%	742.250000	256.750000	408.250000
50%	1526.000000	816.500000	965.500000
75%	3554.250000	3922.000000	1820.250000
max	60869.000000	40827.000000	47943.000000
Number of duplicate rows: 0			

df.corr()

	Channel	Region	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicatessen
Channel	1.000000	0.062028	-0.169172	0.460720	0.608792	-0.202046	0.636026	0.056011
Region	0.062028	1.000000	0.055287	0.032288	0.007696	-0.021044	-0.001483	0.045212
Fresh	-0.169172	0.055287	1.000000	0.100510	-0.011854	0.345881	-0.101953	0.244690
Milk	0.460720	0.032288	0.100510	1.000000	0.728335	0.123994	0.661816	0.406368
Grocery	0.608792	0.007696	-0.011854	0.728335	1.000000	-0.040193	0.924641	0.205497
Frozen	-0.202046	-0.021044	0.345881	0.123994	-0.040193	1.000000	-0.131525	0.390947
Detergents_Paper	0.636026	-0.001483	-0.101953	0.661816	0.924641	-0.131525	1.000000	0.069291
Delicatessen	0.056011	0.045212	0.244690	0.406368	0.205497	0.390947	0.069291	1.000000

```
import seaborn as sns
import matplotlib.pyplot as plt


# boxplots for all features
for column in df.columns:
    plt.figure(figsize=(6, 4))
    sns.boxplot(df[column])
    plt.title(f'Boxplot of {column}')
    plt.show()

# Function to detect outliers
def detect_outliers(dataframe, column):
    Q1 = dataframe[column].quantile(0.25)
    Q3 = dataframe[column].quantile(0.75)
    IQR = Q3 - Q1
    outliers = dataframe[(dataframe[column] < Q1 - 1.5*IQR) | (dataframe[column] > Q3 + 1.5*IQR)]
    return outliers


# number of outliers for each feature
for column in df.columns:
    outliers = detect_outliers(df, column)
    print(f'Number of outliers in {column}: {len(outliers)}')
```




Boxplot of Channel




Boxplot of Region



Boxplot of Fresh




Boxplot of Milk




Boxplot of Grocery







Boxplot of Frozen



Boxplot of Detergents_Paper



Boxplot of Delicatessen




```
def handle_outliers(dataframe, column):
    Q1 = dataframe[column].quantile(0.25)
    Q3 = dataframe[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_limit = Q1 - 1.5*IQR
    upper_limit = Q3 + 1.5*IQR
    dataframe[column] = dataframe[column].apply(lambda x: upper_limit if x > upper_limit else lower_limit if x < lower_limit else x)
```


```
for column in df.columns:
    handle_outliers(df, column)
```

```
for column in df.columns:
    plt.figure(figsize=(6, 4))
    sns.boxplot(df[column])
    plt.title(f'Boxplot of {column}')
    plt.show()
```



Boxplot of Channel




Boxplot of Region




Boxplot of Fresh




Boxplot of Milk




Boxplot of Grocery







Boxplot of Frozen



Boxplot of Detergents_Paper



Boxplot of Delicatessen



```
def detect_outliers(dataframe, column):
    Q1 = dataframe[column].quantile(0.25)
    Q3 = dataframe[column].quantile(0.75)
    IQR = Q3 - Q1
    outliers = dataframe[(dataframe[column] < Q1 - 1.5*IQR) | (dataframe[column] > Q3 + 1.5*IQR)]
    return outliers
```

```
for column in df.columns:
    outliers = detect_outliers(df, column)
    print(f'Number of outliers in {column}: {len(outliers)}')
```

```
Number of outliers in Channel: 0
Number of outliers in Region: 0
Number of outliers in Fresh: 0
Number of outliers in Milk: 0
Number of outliers in Grocery: 0
Number of outliers in Frozen: 0
Number of outliers in Detergents_Paper: 0
Number of outliers in Delicatessen: 0
```


```
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
df_scaled = pd.DataFrame(scaler.fit_transform(df), columns=df.columns)
```

```
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt

# Calculate WCSS for different number of clusters
wcss = []
max_clusters = 15
for i in range(1, max_clusters+1):
    kmeans = KMeans(n_clusters=i, init='k-means++', random_state=42)
    kmeans.fit(df)
    wcss.append(kmeans.inertia_)

# Plot the WCSS values
plt.plot(range(1, max_clusters+1), wcss)
plt.title('The Elbow Method')
plt.xlabel('Number of clusters')
plt.ylabel('WCSS')
plt.grid(True)
plt.show()
```

[illegible]

```
from sklearn.cluster import KMeans

# Build the model
kmeans = KMeans(n_clusters=4, init='k-means++', random_state=42)
kmeans.fit(df)

# Get cluster labels
```

```
cluster_labels = kmeans.labels_  
  
# Add cluster labels to your original dataframe  
df['Cluster'] = cluster_labels
```

/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' warnings.warn(

```
df['Cluster'] = kmeans.labels_  
  
# Check the size of each cluster  
print("Cluster Sizes:\n", df['Cluster'].value_counts())  
  
# Check the characteristics of each cluster  
for i in range(4):  
    print("\nCluster ", i)  
    print(df[df['Cluster'] == i].describe())
```

Cluster Sizes:
3 176
0 112
1 94
2 58
Name: Cluster, dtype: int64

Cluster 0

	Channel	Region	Fresh	Milk	Grocery \
count	112.000000	112.000000	112.000000	112.000000	112.000000
mean	1.214286	2.535714	16051.205357	3135.813616	4211.589286
std	0.412170	0.781873	3763.633078	2524.464860	3150.441587
min	1.000000	1.000000	10379.000000	134.000000	3.000000
25%	1.000000	2.000000	12419.750000	1283.500000	1970.500000
50%	1.000000	3.000000	16195.000000	2252.000000	3203.000000
75%	1.000000	3.000000	18830.250000	4537.000000	5700.250000
max	2.000000	3.000000	24929.000000	15676.125000	14982.000000

	Frozen	Detergents_Paper	Delicatessen	Cluster
count	112.000000	112.000000	112.000000	112.0
mean	2988.859375	994.785714	1229.573661	0.0
std	2531.352938	1245.589613	963.527882	0.0
min	118.000000	3.000000	51.000000	0.0
25%	1018.750000	188.500000	514.250000	0.0
50%	2157.500000	456.500000	879.000000	0.0
75%	4276.000000	1404.000000	1804.500000	0.0
max	7772.250000	6707.000000	3938.250000	0.0

Cluster 1

	Channel	Region	Fresh	Milk	Grocery \
count	94.000000	94.000000	94.000000	94.000000	94.000000
mean	1.893617	2.489362	5331.893617	10454.450798	17196.140957
std	0.309980	0.799794	5111.448153	3937.245330	4905.345002
min	1.000000	1.000000	18.000000	1266.000000	8852.000000
25%	2.000000	2.000000	1409.500000	7576.000000	12563.250000
50%	2.000000	3.000000	4047.000000	10601.000000	16596.000000
75%	2.000000	3.000000	7870.500000	14316.500000	22288.500000
max	2.000000	3.000000	22925.000000	15676.125000	23409.875000

	Frozen	Detergents_Paper	Delicatessen	Cluster
count	94.000000	94.000000	94.000000	94.0
mean	1496.428191	6936.898936	1547.364362	1.0
std	1538.882840	2383.035957	1176.131062	0.0
min	25.000000	241.000000	3.000000	1.0
25%	438.500000	5274.250000	680.000000	1.0
50%	973.000000	6931.500000	1366.500000	1.0
75%	1900.000000	9419.875000	2157.750000	1.0
max	7772.250000	9419.875000	3938.250000	1.0

Cluster 2

	Channel	Region	Fresh	Milk	Grocery \
count	58.000000	58.000000	58.000000	58.000000	58.000000
mean	1.172414	2.655172	32136.810345	5973.515086	7309.012931
std	0.381039	0.714554	5122.024937	4808.223223	5915.174661
min	1.000000	1.000000	22647.000000	286.000000	471.000000
25%	1.000000	3.000000	27207.500000	2393.000000	2726.250000
50%	1.000000	3.000000	31664.000000	4347.000000	5259.500000
75%	1.000000	3.000000	37642.750000	7829.500000	9344.000000

```
from sklearn.decomposition import PCA  
import matplotlib.pyplot as plt  
  
# Apply PCA and fit the features selected  
pca = PCA(n_components=2)  
principalComponents = pca.fit_transform(df.drop('Cluster', axis=1))  
  
# Create a DataFrame with the two components  
PCA_components = pd.DataFrame(principalComponents, columns=['Principal Component 1', 'Principal Component 2'])
```

```
# Concatenate the clusters labels to the DataFrame
```

```
PCA_components['Cluster'] = df['Cluster']
```

```
# Plot the clustered dataset
```

```
plt.figure(figsize=(8,6))
```

```
plt.scatter(PCA_components['Principal Component 1'], PCA_components['Principal Component 2'], c=PCA_components['Cluster'])
```

```
plt.title('Clusters in PCA 2D Space')
```

```
plt.xlabel('Principal Component 1')
```

```
plt.ylabel('Principal Component 2')
```

```
plt.colorbar(label='Cluster')
```

```
plt.show()
```

